

# **Reading Material for System Call Implementation in Linux Operating System [Including Kernel Recompilation]**

**As Part of  
Operating Systems [CS F372] Course  
Semester I, 2019 - 2020**



**BITS Pilani**  
K K Birla Goa Campus

**Prepared by  
HARSHA V, VEERENDRA VIDYADHAR RISBUD,  
PALLAVI VARSHNEY AND KARUNA GREWAL**

**Edited by  
BIJU K RAVEENDRAN, Dept. of CS & IS  
BIRLA INSTITUTE OF TECHNOLOGY AND  
SCIENCE, PILANI –K K BIRLA GOA CAMPUS**

# INDEX

<b>Sec. No.</b>	<b>Section Title</b>	<b>Page No.</b>
<b>1.</b>	<b>Basic Preparation for Recompilation</b>	<b>3</b>
<b>2.</b>	<b>Recompilation of Linux Kernel with / without Modification(s)</b>	<b>5</b>
<b>3.</b>	<b>Implementation of New System Call [New System Call to Add Two Positive Integers]</b>	<b>9</b>
<b>4.</b>	<b>Implementation of User Space Programs</b>	<b>15</b>

# Basic Preparation for Recompilation

Step #1. Download the kernel source 5.2.9 from <ftp://10.1.9.224/Linux/Kernels/> or from <https://www.kernel.org>

Step #2: Open a terminal and login to super-user by

```
$ sudo su
<Enter root password here>
```

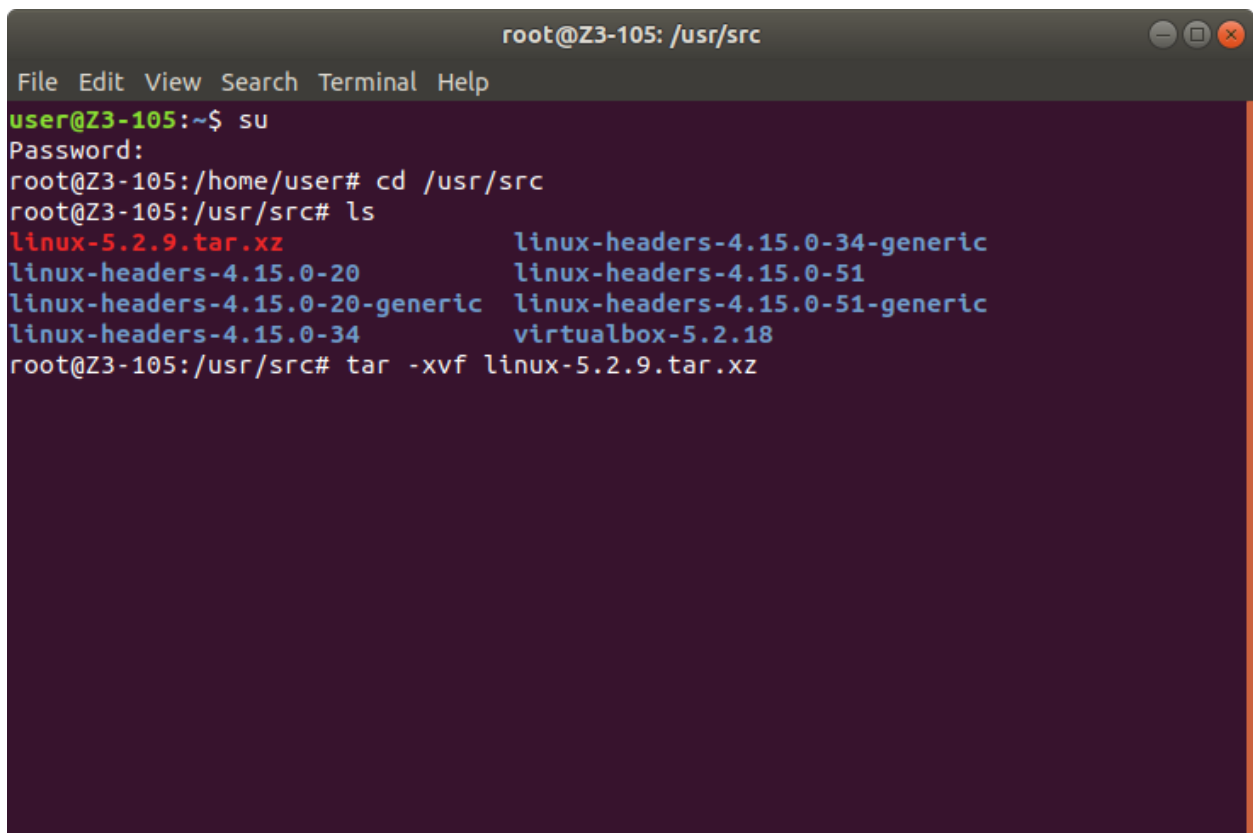
Step #3: Place the tar.xz file in /usr/src/ directory

Step #4: Set the present working directory as /usr/src/ by

```
$ cd /usr/src/
```

Step #5: Untar the *linux-5.2.9.tar.xz* file by

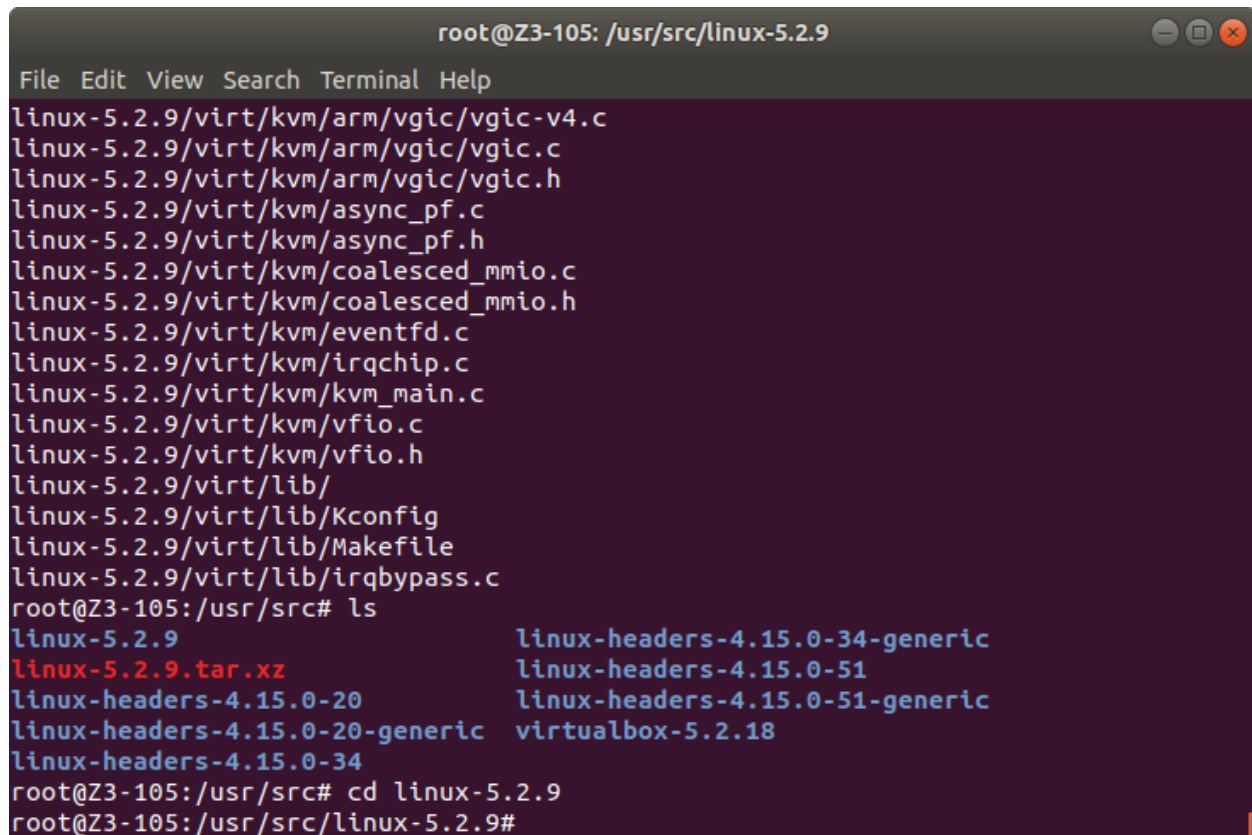
```
$ tar -xvf linux-5.2.9.tar.xz
```

A terminal window titled 'root@Z3-105: /usr/src' with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal shows the following commands and output:

```
user@Z3-105:~$ su
Password:
root@Z3-105:/home/user# cd /usr/src
root@Z3-105:/usr/src# ls
linux-5.2.9.tar.xz          linux-headers-4.15.0-34-generic
linux-headers-4.15.0-20    linux-headers-4.15.0-51
linux-headers-4.15.0-20-generic  linux-headers-4.15.0-51-generic
linux-headers-4.15.0-34      virtualbox-5.2.18
root@Z3-105:/usr/src# tar -xvf linux-5.2.9.tar.xz
```

Step #6: Set the present working directory as linux-5.2.9

**\$ cd linux-5.2.9/**



```
root@Z3-105: /usr/src/linux-5.2.9
File Edit View Search Terminal Help
linux-5.2.9/virt/kvm/arm/vgic/vgic-v4.c
linux-5.2.9/virt/kvm/arm/vgic/vgic.c
linux-5.2.9/virt/kvm/arm/vgic/vgic.h
linux-5.2.9/virt/kvm/async_pf.c
linux-5.2.9/virt/kvm/async_pf.h
linux-5.2.9/virt/kvm/coalesced_mmio.c
linux-5.2.9/virt/kvm/coalesced_mmio.h
linux-5.2.9/virt/kvm/eventfd.c
linux-5.2.9/virt/kvm/irqchip.c
linux-5.2.9/virt/kvm/kvm_main.c
linux-5.2.9/virt/kvm/vfio.c
linux-5.2.9/virt/kvm/vfio.h
linux-5.2.9/virt/lib/
linux-5.2.9/virt/lib/Kconfig
linux-5.2.9/virt/lib/Makefile
linux-5.2.9/virt/lib/irqbypass.c
root@Z3-105:/usr/src# ls
linux-5.2.9                linux-headers-4.15.0-34-generic
linux-5.2.9.tar.xz         linux-headers-4.15.0-51
linux-headers-4.15.0-20    linux-headers-4.15.0-51-generic
linux-headers-4.15.0-20-generic  virtualbox-5.2.18
linux-headers-4.15.0-34
root@Z3-105:/usr/src# cd linux-5.2.9
root@Z3-105:/usr/src/linux-5.2.9#
```

# Recompilation of Linux Kernel with / without Modification(s)

## Step #1: Reconfiguration of the Kernel

The Linux Kernel is extraordinarily configurable; you can enable and disable many of its features, as well as set build parameters. Some of the widely used options are: menuconfig, xconfig, gconfig, oldconfig, defconfig etc.

**Dependencies you may require to install: flex, bison, libssl-dev**

**\$apt install flex bison libssl-dev                      or**

**\$apt-get install <package name>      E.g.: \$apt-get install flex**

**\$ make menuconfig**

<Text based color menus, radio lists & dialogs. This option is also useful on remote server if you want to compile kernel remotely.>

**\$ make xconfig**

<X windows (Qt) based configuration tool, works best under KDE Desktop.>

**\$ make gconfig**

<X windows (Gtk) based configuration tool, works best under Gnome Desktop.>

**\$ make oldconfig**

<Reads the existing config file and prompts the user options in the current kernel source that are not found in the file>

**\$ make defconfig [Use this for reconfiguration option for this assignment]**

<Creates a default config file for the kernel delineating all the necessary modules to be installed into the kernel>

```
root@Z3-105: /usr/src/linux-5.2.9
File Edit View Search Terminal Help
linux-5.2.9          linux-headers-4.15.0-34-generic
linux-5.2.9.tar.xz   linux-headers-4.15.0-51
linux-headers-4.15.0-20  linux-headers-4.15.0-51-generic
linux-headers-4.15.0-20-generic virtualbox-5.2.18
linux-headers-4.15.0-34
root@Z3-105:/usr/src# cd linux-5.2.9
root@Z3-105:/usr/src/linux-5.2.9# make defconfig
HOSTCC  scripts/basic/fixdep
HOSTCC  scripts/kconfig/conf.o
HOSTCC  scripts/kconfig/confdata.o
HOSTCC  scripts/kconfig/expr.o
LEX      scripts/kconfig/lexer.lex.c
YACC     scripts/kconfig/parser.tab.h
HOSTCC  scripts/kconfig/lexer.lex.o
YACC     scripts/kconfig/parser.tab.c
HOSTCC  scripts/kconfig/parser.tab.o
HOSTCC  scripts/kconfig/preprocess.o
HOSTCC  scripts/kconfig/symbol.o
HOSTLD  scripts/kconfig/conf
*** Default configuration is based on 'x86_64_defconfig'
#
# configuration written to .config
#
root@Z3-105:/usr/src/linux-5.2.9#
```

## Step #2: Preliminary Recompilation of the Kernel

Execute make to compile the kernel.

\$ make

```
root@Z3-105: /usr/src/linux-5.2.9
File Edit View Search Terminal Help
root@Z3-105:/usr/src/linux-5.2.9# make
```

```
root@Z3-105: /usr/src/linux-5.2.9
File Edit View Search Terminal Help
LD [M] drivers/thermal/intel/x86_pkg_temp_thermal.ko
CC fs/efivarfs/efivarfs.mod.o
LD [M] fs/efivarfs/efivarfs.ko
CC net/ipv4/netfilter/iptable_nat.mod.o
LD [M] net/ipv4/netfilter/iptable_nat.ko
CC net/ipv4/netfilter/nf_log_arp.mod.o
LD [M] net/ipv4/netfilter/nf_log_arp.ko
CC net/ipv4/netfilter/nf_log_ipv4.mod.o
LD [M] net/ipv4/netfilter/nf_log_ipv4.ko
CC net/ipv6/netfilter/nf_log_ipv6.mod.o
LD [M] net/ipv6/netfilter/nf_log_ipv6.ko
CC net/netfilter/nf_log_common.mod.o
LD [M] net/netfilter/nf_log_common.ko
CC net/netfilter/xt_LOG.mod.o
LD [M] net/netfilter/xt_LOG.ko
CC net/netfilter/xt_MASQUERADE.mod.o
LD [M] net/netfilter/xt_MASQUERADE.ko
CC net/netfilter/xt_addrtype.mod.o
LD [M] net/netfilter/xt_addrtype.ko
CC net/netfilter/xt_mark.mod.o
LD [M] net/netfilter/xt_mark.ko
CC net/netfilter/xt_nat.mod.o
LD [M] net/netfilter/xt_nat.ko
root@Z3-105: /usr/src/linux-5.2.9#
```

### Step #3: Recompilation of the Kernel module, update initramfs and grub

Execute `make modules_install` & `make modules_install install` to compile the modules and update the initramfs and grub.

**\$ `make modules_install && make modules_install install`**

In addition to installing the bzImage it even runs the following commands

`update-initramfs -c -k linux-5.2.9`

`update-grub`

```
root@Z3-105: /usr/src/linux-5.2.9
File Edit View Search Terminal Help
root@Z3-105: /usr/src/linux-5.2.9# make modules_install && make modules_install install
```

```
root@Z3-105: /usr/src/linux-5.2.9
File Edit View Search Terminal Help
INSTALL net/ipv4/netfilter/iptable_nat.ko
INSTALL net/ipv4/netfilter/nf_log_arp.ko
INSTALL net/ipv4/netfilter/nf_log_ipv4.ko
INSTALL net/ipv6/netfilter/nf_log_ipv6.ko
INSTALL net/netfilter/nf_log_common.ko
INSTALL net/netfilter/xt_LOG.ko
INSTALL net/netfilter/xt_MASQUERADE.ko
INSTALL net/netfilter/xt_addrtype.ko
INSTALL net/netfilter/xt_mark.ko
INSTALL net/netfilter/xt_nat.ko
DEPMOD 5.2.9
sh ./arch/x86/boot/install.sh 5.2.9 arch/x86/boot/bzImage \
    System.map "/boot"
run-parts: executing /etc/kernel/postinst.d/apt-auto-removal 5.2.9 /boot/vmlinuz-5.2.9
run-parts: executing /etc/kernel/postinst.d/dkms 5.2.9 /boot/vmlinuz-5.2.9
ERROR (dkms apport): kernel package linux-headers-5.2.9 is not supported
Error! Bad return status for module build on kernel: 5.2.9 (x86_64)
Consult /var/lib/dkms/virtualbox/5.2.18/build/make.log for more information.
run-parts: executing /etc/kernel/postinst.d/initramfs-tools 5.2.9 /boot/vmlinuz-5.2.9
update-initramfs: Generating /boot/initrd.img-5.2.9
run-parts: executing /etc/kernel/postinst.d/unattended-upgrades 5.2.9 /boot/vmlinuz-5.2.9
run-parts: executing /etc/kernel/postinst.d/update-notifier 5.2.9 /boot/vmlinuz-5.2.9
run-parts: executing /etc/kernel/postinst.d/zz-update-grub 5.2.9 /boot/vmlinuz-5.2.9
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-5.2.9
Found initrd image: /boot/initrd.img-5.2.9
Found linux image: /boot/vmlinuz-4.15.0-51-generic
Found initrd image: /boot/initrd.img-4.15.0-51-generic
Found linux image: /boot/vmlinuz-4.15.0-34-generic
Found initrd image: /boot/initrd.img-4.15.0-34-generic
Found linux image: /boot/vmlinuz-4.15.0-20-generic
Found initrd image: /boot/initrd.img-4.15.0-20-generic
Found memtest86+ image: /memtest86+.elf
Found memtest86+ image: /memtest86+.bin
Found Windows 10 on /dev/sda1
done
root@Z3-105: /usr/src/linux-5.2.9#
```

**Now that the kernel has been recompiled, reboot the system and boot into this kernel from the grub <Select advanced ubuntu tab followed by the New kernel>**

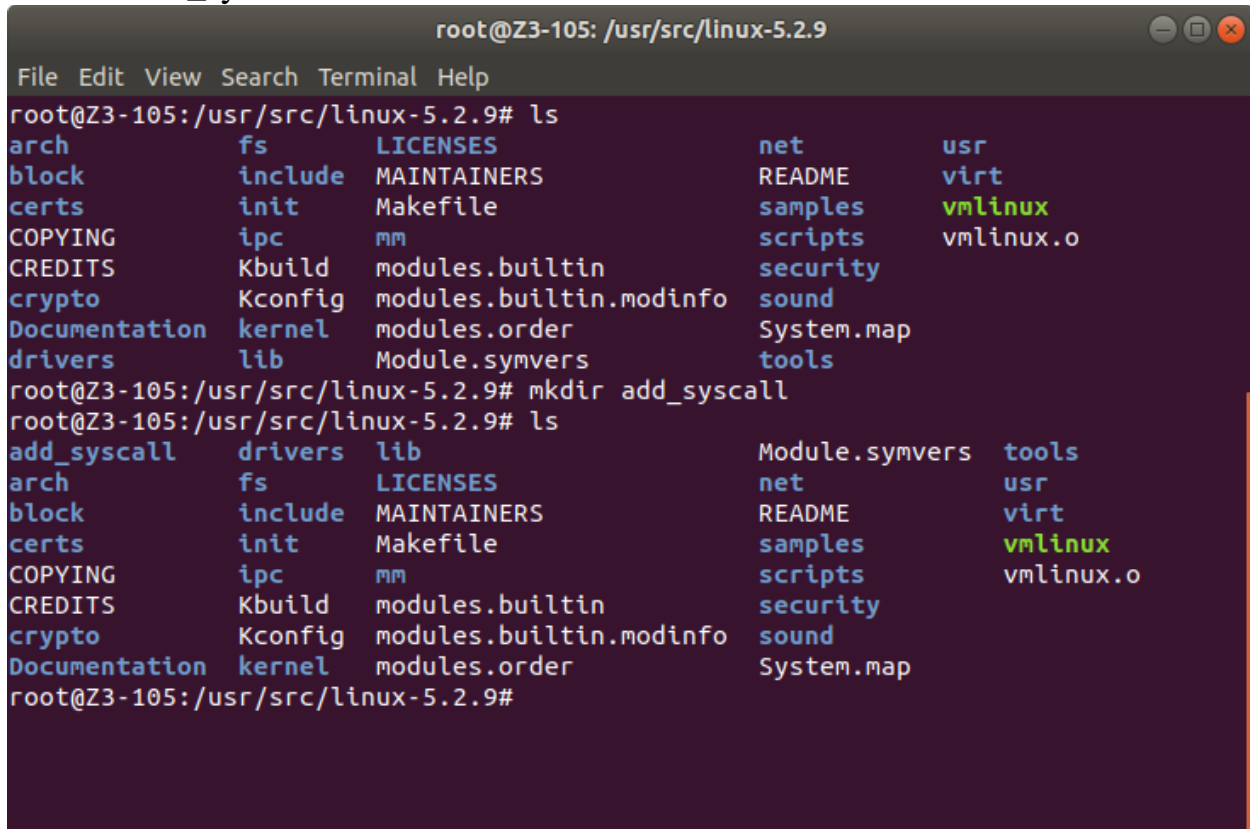


# Implementation of New System Call [New System Call to Add 2 Positive Integers]

Step #1: Create a directory under /usr/src/linux-5.2.9/

Create a directory named add\_syscall under /usr/src/linux-5.2.9/

\$ mkdir add\_syscall



```
root@Z3-105: /usr/src/linux-5.2.9
File Edit View Search Terminal Help
root@Z3-105:/usr/src/linux-5.2.9# ls
arch          fs            LICENSES      net           usr
block         include      MAINTAINERS   README       virt
certs         init         Makefile      samples      vmlinux
COPYING       ipc          mm            scripts       vmlinux.o
CREDITS       Kbuild       modules.builtin  security
crypto        Kconfig     modules.builtin.modinfo  sound
Documentation kernel       modules.order  System.map
drivers       lib          Module.symvers tools
root@Z3-105:/usr/src/linux-5.2.9# mkdir add_syscall
root@Z3-105:/usr/src/linux-5.2.9# ls
add_syscall   drivers       lib           Module.symvers  tools
arch          fs            LICENSES      net             usr
block         include      MAINTAINERS   README          virt
certs         init         Makefile      samples         vmlinux
COPYING       ipc          mm            scripts          vmlinux.o
CREDITS       Kbuild       modules.builtin  security
crypto        Kconfig     modules.builtin.modinfo  sound
Documentation kernel       modules.order  System.map
root@Z3-105:/usr/src/linux-5.2.9#
```

Step #2: Create the following files under add\_syscall directory

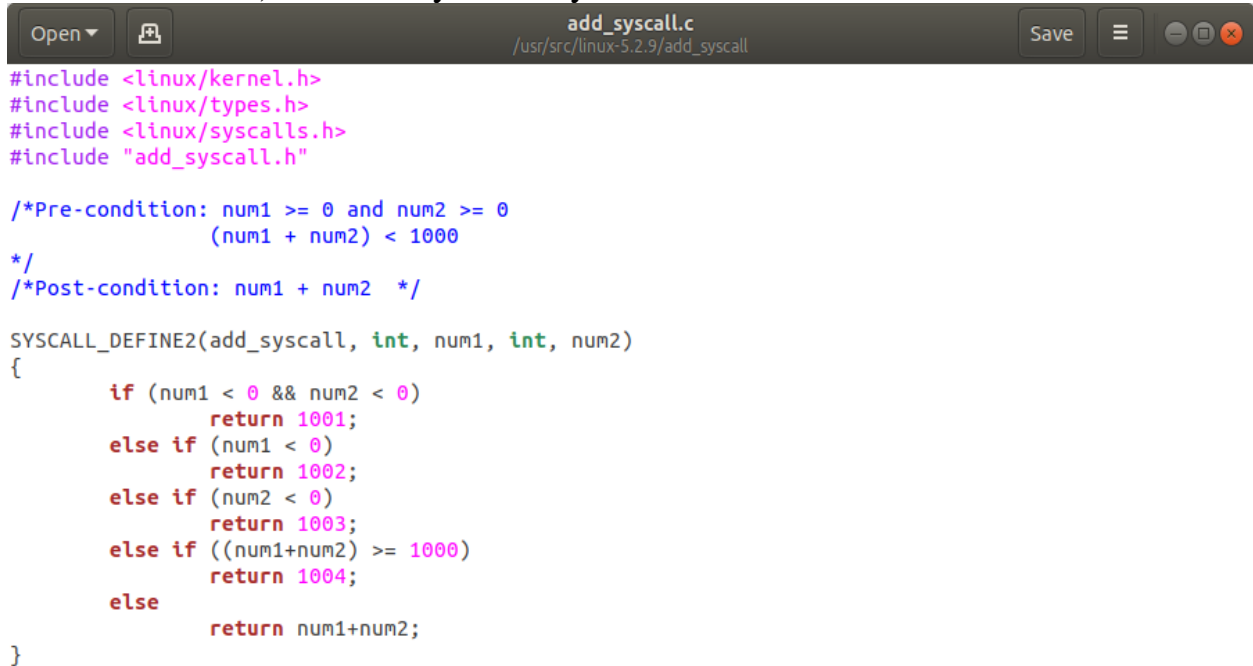
1. add\_syscall.c
2. add\_syscall.h
3. Makefile

Contents of **add\_syscall.c**

SYSCALL\_DEFINE $n$ () macros are the standard way for kernel code to define a system call, where the  $n$  suffix indicates the argument count.

The first argument to the macro is the name of the system call (without sys\_ prepended to it). The remaining arguments are pairs of type and name for the parameters.

The definitions of these **SYSCALL\_DEFINE...** macros are in **#include <linux/syscalls.h>**. Hence, the .c file in which you code the body of your syscall's service routine must **#include <linux/syscalls.h>** It has within { ... } (after your SYSCALL\_DEFINE...(...) ) the code (you will write!) of the body of the syscall to be run.



```

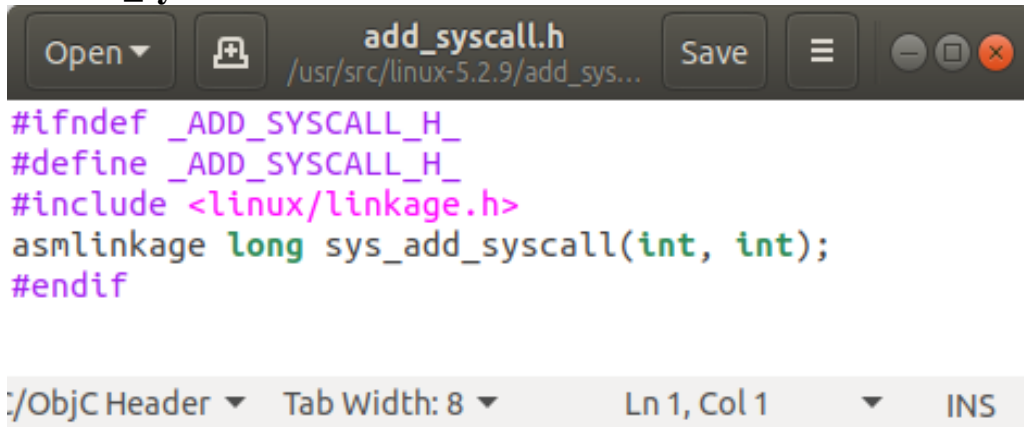
#include <linux/kernel.h>
#include <linux/types.h>
#include <linux/syscalls.h>
#include "add_syscall.h"

/*Pre-condition: num1 >= 0 and num2 >= 0
               (num1 + num2) < 1000
*/
/*Post-condition: num1 + num2 */

SYSCALL_DEFINE2(add_syscall, int, num1, int, num2)
{
    if (num1 < 0 && num2 < 0)
        return 1001;
    else if (num1 < 0)
        return 1002;
    else if (num2 < 0)
        return 1003;
    else if ((num1+num2) >= 1000)
        return 1004;
    else
        return num1+num2;
}

```

Content of **add\_syscall.h**



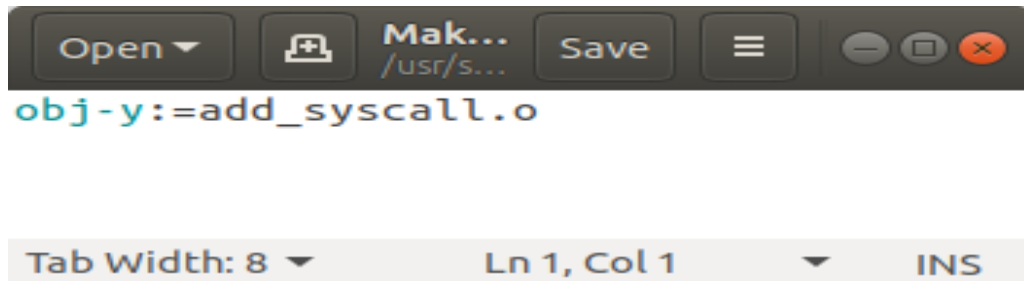
```

#ifndef _ADD_SYSCALL_H_
#define _ADD_SYSCALL_H_
#include <linux/linkage.h>
asmlinkage long sys_add_syscall(int, int);
#endif

```

:/ObjC Header ▾ Tab Width: 8 ▾ Ln 1, Col 1 ▾ INS

Content of **Makefile**



```

obj-y:=add_syscall.o

```

Tab Width: 8 ▾ Ln 1, Col 1 ▾ INS

## Step #3: Modify the following files

1. `/usr/src/linux-5.2.9/Makefile`
2. `/usr/src/linux-5.2.9/arch/x86/entry/syscalls/syscall_64.tbl`
3. `/usr/src/linux-5.2.9/include/asm-generic/syscalls.h`
4. `/usr/src/linux-5.2.9/include/linux/syscalls.h`

### 3.1: Modify `/usr/src/linux-5.2.9/Makefile`:

<Update the following line in Makefile>

`core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/`

<to the following by adding `add_syscall/` in the end>

`core-y += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ add_syscall/`



```
HOST_LIBELF_LIBS = $(shell pkg-config libelf --libs 2>/dev/null || echo -lelf)

ifndef CONFIG_STACK_VALIDATION
    has_libelf := $(call try-run,\
        echo "int main() {}" | $(HOSTCC) -xc -o /dev/null $(HOST_LIBELF_LIBS) -,1,0)
    ifeq ($(has_libelf),1)
        objtool_target := tools/objtool FORCE
    else
        SKIP_STACK_VALIDATION := 1
        export SKIP_STACK_VALIDATION
    endif
endif

PHONY += prepare0

ifeq ($(KBUILD_EXTMOD),)
core-y      += kernel/ certs/ mm/ fs/ ipc/ security/ crypto/ block/ add_syscall/

vmlinux-dirs := $(patsubst %/,%, $(filter %/, $(init-y) $(init-m) \
    $(core-y) $(core-m) $(drivers-y) $(drivers-m) \
    $(net-y) $(net-m) $(libs-y) $(libs-m) $(virt-y)))

vmlinux-alldirs := $(sort $(vmlinux-dirs) Documentation \
    $(patsubst %/,%, $(filter %/, $(init-) $(core-) \
    $(drivers-) $(net-) $(libs-) $(virt-))))

init-y      := $(patsubst %/, %/built-in.a, $(init-y))
core-y      := $(patsubst %/, %/built-in.a, $(core-y))
drivers-y   := $(patsubst %/, %/built-in.a, $(drivers-y))
net-y       := $(patsubst %/, %/built-in.a, $(net-y))
libs-y1     := $(patsubst %/, %/lib.a, $(libs-y))
libs-y2     := $(patsubst %/, %/built-in.a, $(filter-out %.a, $(libs-y)))
virt-y      := $(patsubst %/, %/built-in.a, $(virt-y))

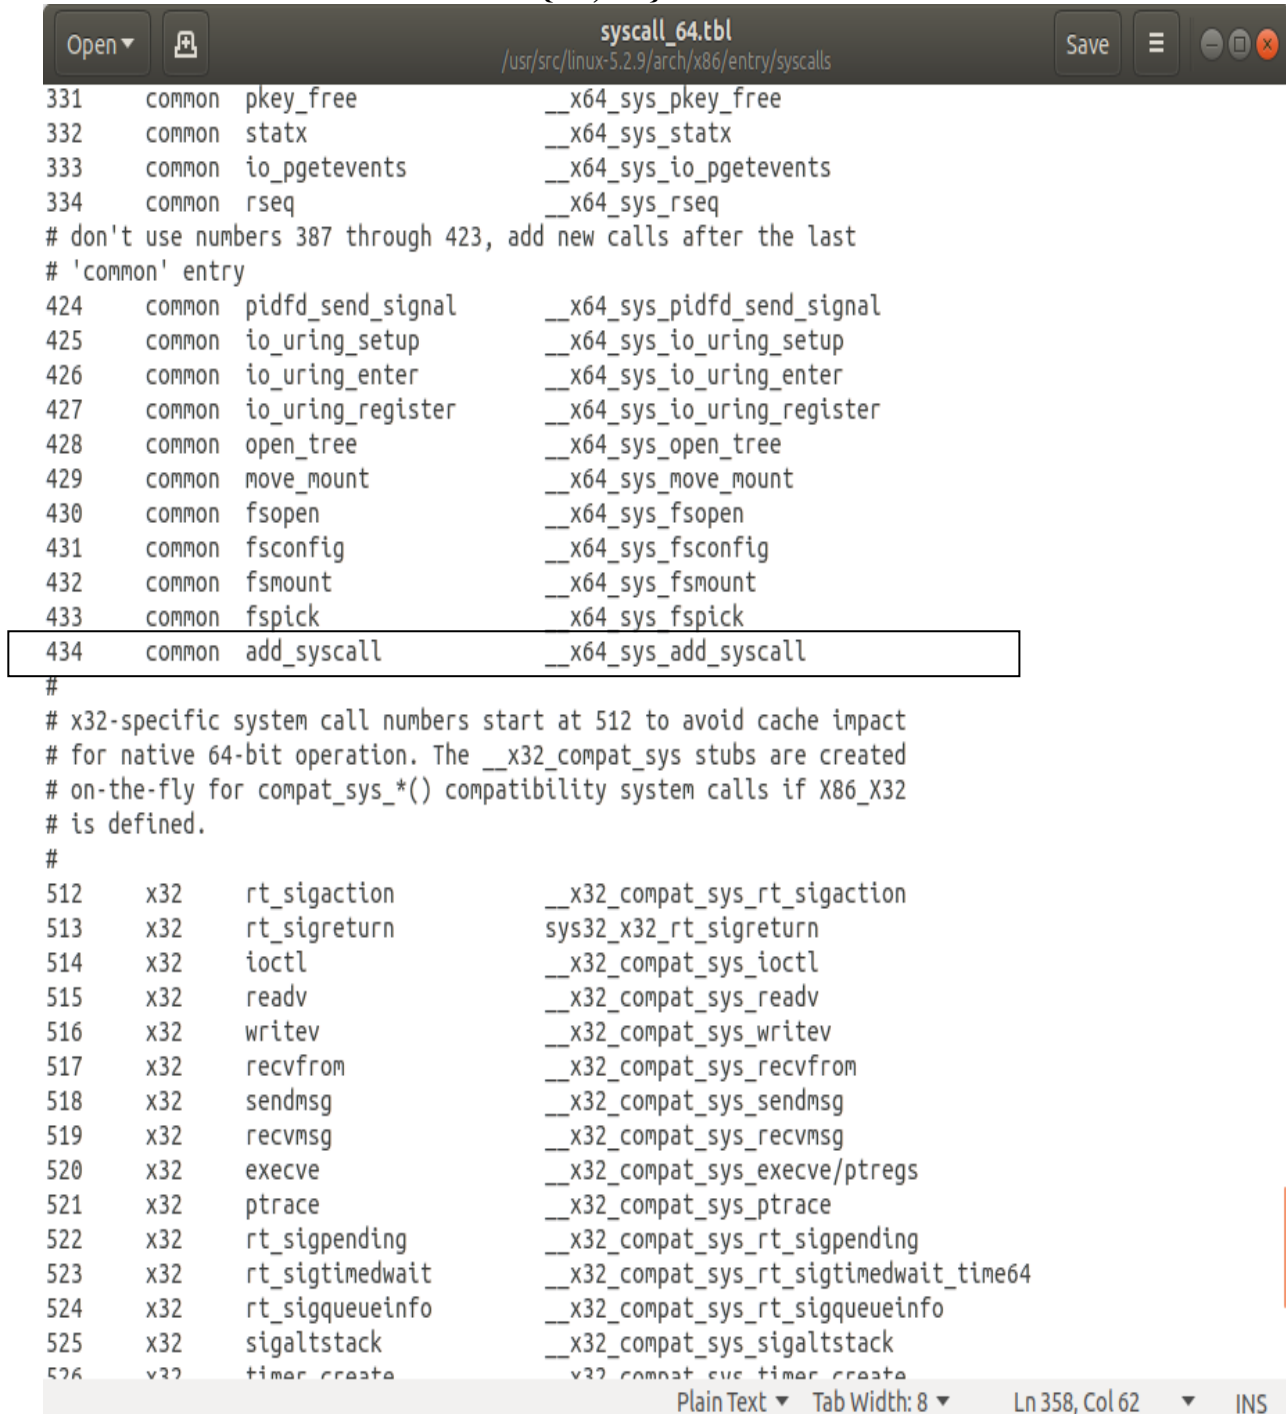
# Externally visible symbols (used by link-vmlinux.sh)
export KBUILD_VMLINUX_OBJS := $(head-y) $(init-y) $(core-y) $(libs-y2) \
    $(drivers-y) $(net-y) $(virt-y)
```

### 3.2: Modify /usr/src/linux-5.2.9/arch/x86/entry/syscalls/syscall\_64.tbl:

Update the file: /arch/x86/entry/syscalls/syscall\_64.tbl to add the new syscall at the next available system call number in the common list of syscalls like:

**434 common add\_syscall \_\_x64\_sys\_add\_syscall**

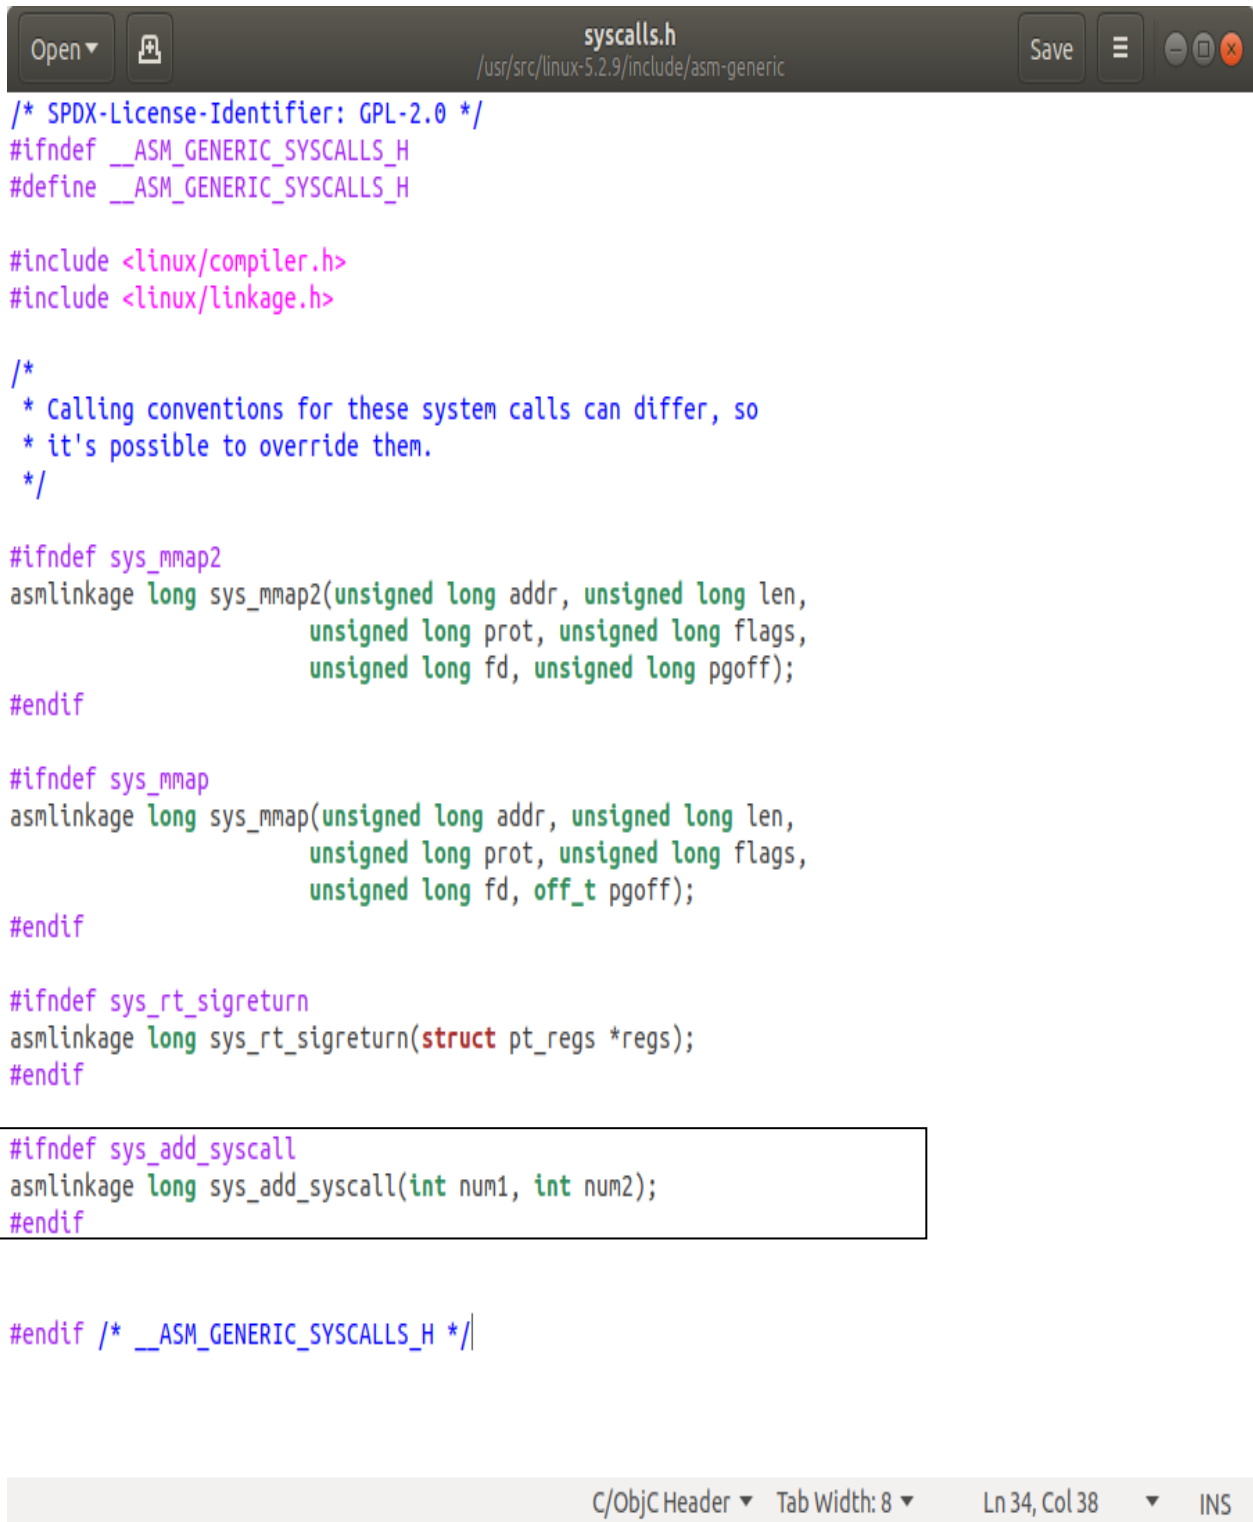
**Here sys\_add\_syscall is the entry point for the system call add\_syscall and it will be common across the x86-{64, 32} bit architectures.**



```
Open ▾  syscall_64.tbl  Save  ≡  -  □  ×
/usr/src/linux-5.2.9/arch/x86/entry/syscalls
331    common    pkey_free        __x64_sys_pkey_free
332    common    statx            __x64_sys_statx
333    common    io_pgetevents    __x64_sys_io_pgetevents
334    common    rseq             __x64_sys_rseq
# don't use numbers 387 through 423, add new calls after the last
# 'common' entry
424    common    pidfd_send_signal  __x64_sys_pidfd_send_signal
425    common    io_uring_setup    __x64_sys_io_uring_setup
426    common    io_uring_enter    __x64_sys_io_uring_enter
427    common    io_uring_register  __x64_sys_io_uring_register
428    common    open_tree        __x64_sys_open_tree
429    common    move_mount        __x64_sys_move_mount
430    common    fsopen           __x64_sys_fsopen
431    common    fsconfig         __x64_sys_fsconfig
432    common    fsmount          __x64_sys_fsmount
433    common    fspick           x64 sys fspick
434    common    add_syscall       __x64_sys_add_syscall
#
# x32-specific system call numbers start at 512 to avoid cache impact
# for native 64-bit operation. The __x32_compat_sys stubs are created
# on-the-fly for compat_sys_*( ) compatibility system calls if X86_X32
# is defined.
#
512    x32       rt_sigaction     __x32_compat_sys_rt_sigaction
513    x32       rt_sigreturn     sys32_x32_rt_sigreturn
514    x32       ioctl            __x32_compat_sys_ioctl
515    x32       readv            __x32_compat_sys_readv
516    x32       writev           __x32_compat_sys_writev
517    x32       recvfrom         __x32_compat_sys_recvfrom
518    x32       sendmsg          __x32_compat_sys_sendmsg
519    x32       recvmmsg         __x32_compat_sys_recvmmsg
520    x32       execve           __x32_compat_sys_execve/ptregs
521    x32       ptrace           __x32_compat_sys_ptrace
522    x32       rt_sigpending     __x32_compat_sys_rt_sigpending
523    x32       rt_sigtimedwait   __x32_compat_sys_rt_sigtimedwait_time64
524    x32       rt_sigqueueinfo   __x32_compat_sys_rt_sigqueueinfo
525    x32       sigaltstack       __x32_compat_sys_sigaltstack
526    x32       timer_create      x32 compat sys timer create
Plain Text ▾  Tab Width: 8 ▾  Ln 358, Col 62  ▾  INS
```

**This table is read by scripts and used to generate some of the boilerplate code**

### 3.3: Modify /usr/src/linux-5.2.9/include/asm-generic/syscalls.h:



```
syscalls.h
/usr/src/linux-5.2.9/include/asm-generic

/* SPDX-License-Identifier: GPL-2.0 */
#ifndef __ASM_GENERIC_SYSCALLS_H
#define __ASM_GENERIC_SYSCALLS_H

#include <linux/compiler.h>
#include <linux/linkage.h>

/*
 * Calling conventions for these system calls can differ, so
 * it's possible to override them.
 */

#ifndef sys_mmap2
asmlinkage long sys_mmap2(unsigned long addr, unsigned long len,
                          unsigned long prot, unsigned long flags,
                          unsigned long fd, unsigned long pgoff);
#endif

#ifndef sys_mmap
asmlinkage long sys_mmap(unsigned long addr, unsigned long len,
                          unsigned long prot, unsigned long flags,
                          unsigned long fd, off_t pgoff);
#endif

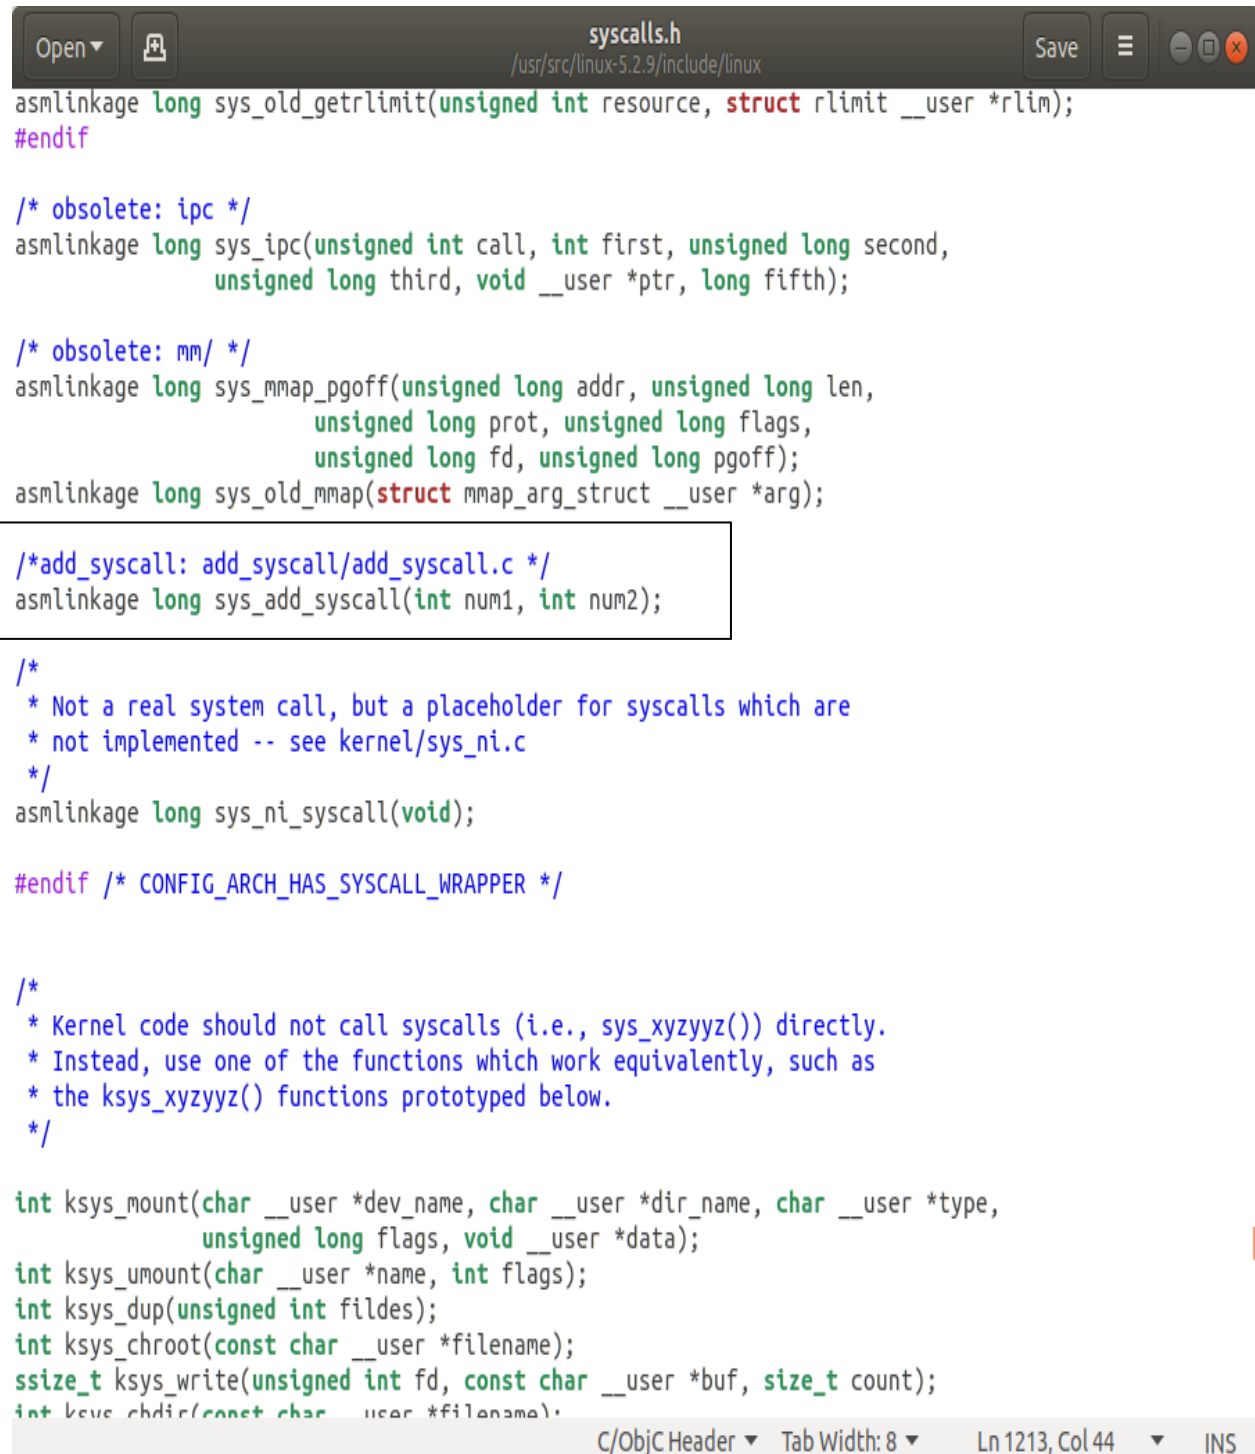
#ifndef sys_rt_sigreturn
asmlinkage long sys_rt_sigreturn(struct pt_regs *regs);
#endif

#ifndef sys_add_syscall
asmlinkage long sys_add_syscall(int num1, int num2);
#endif

#endif /* __ASM_GENERIC_SYSCALLS_H */

C/ObjC Header ▾ Tab Width: 8 ▾ Ln 34, Col 38 ▾ INS
```

### 3.4: Modify /usr/src/linux-5.2.9/include/linux/syscalls.h:



```
asmlinkage long sys_old_getrlimit(unsigned int resource, struct rlimit __user *rlim);
#endif

/* obsolete: ipc */
asmlinkage long sys_ipc(unsigned int call, int first, unsigned long second,
                        unsigned long third, void __user *ptr, long fifth);

/* obsolete: mm/ */
asmlinkage long sys_mmap_pgoff(unsigned long addr, unsigned long len,
                               unsigned long prot, unsigned long flags,
                               unsigned long fd, unsigned long pgoff);
asmlinkage long sys_old_mmap(struct mmap_arg_struct __user *arg);

/*add_syscall: add_syscall/add_syscall.c */
asmlinkage long sys_add_syscall(int num1, int num2);

/*
 * Not a real system call, but a placeholder for syscalls which are
 * not implemented -- see kernel/sys_ni.c
 */
asmlinkage long sys_ni_syscall(void);

#endif /* CONFIG_ARCH_HAS_SYSCALL_WRAPPER */

/*
 * Kernel code should not call syscalls (i.e., sys_xyzzyz()) directly.
 * Instead, use one of the functions which work equivalently, such as
 * the ksys_xyzzyz() functions prototyped below.
 */

int ksys_mount(char __user *dev_name, char __user *dir_name, char __user *type,
               unsigned long flags, void __user *data);
int ksys_umount(char __user *name, int flags);
int ksys_dup(unsigned int fildes);
int ksys_chroot(const char __user *filename);
ssize_t ksys_write(unsigned int fd, const char __user *buf, size_t count);
int ksys_chdir(const char __user *filename);
```

Recompile the Kernel [Follow section#2]to get all the changes reflected. Reboot the system and boot into this kernel from the grub <Select advanced ubuntu tab followed by the New kernel>

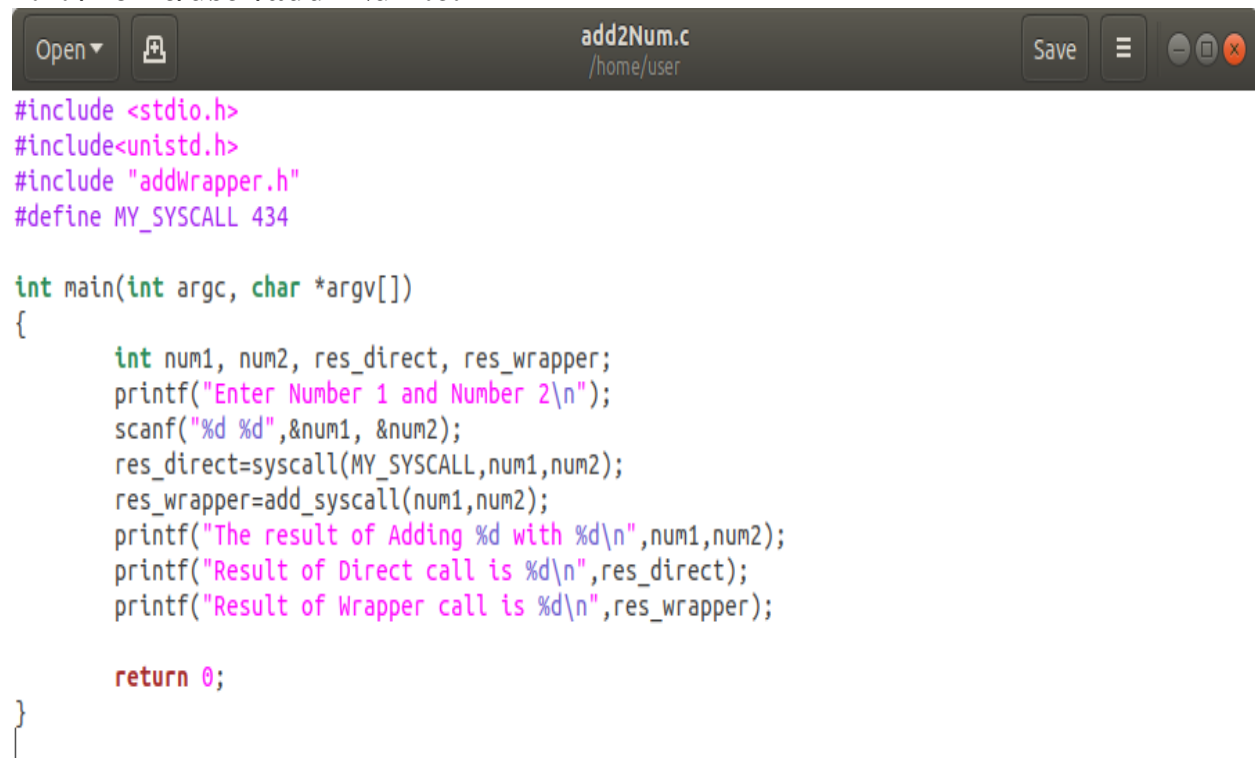
# Implementation of User Space Programs

1. /home/user/add2Num.c

2. /home/user/addWrapper.h

The C user library wraps most system calls for us. This avoids triggering interrupts directly. The user space .c file provides two mechanisms of calling a system call (A) directly using the *syscall()* function with the help of system call number [GNU C library provides this for us] and (B) with the help of a Wrapper where the end user never need to remember the system call number.

## 1.1: /home/user/add2Num.c:

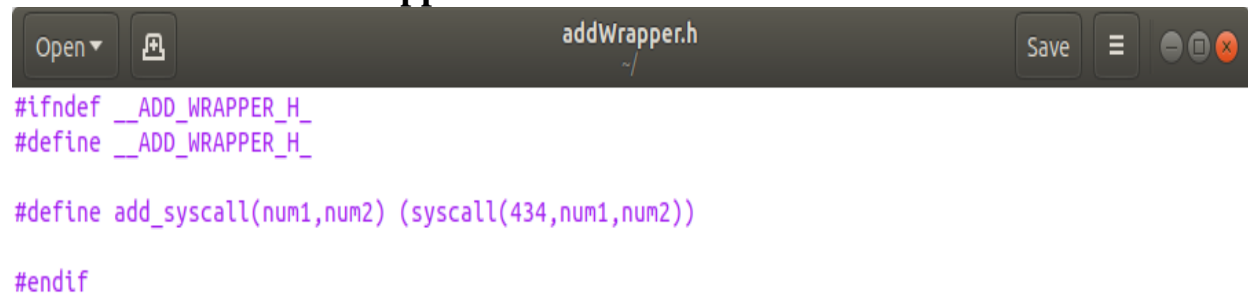


```
#include <stdio.h>
#include <unistd.h>
#include "addWrapper.h"
#define MY_SYSCALL 434

int main(int argc, char *argv[])
{
    int num1, num2, res_direct, res_wrapper;
    printf("Enter Number 1 and Number 2\n");
    scanf("%d %d", &num1, &num2);
    res_direct=syscall(MY_SYSCALL,num1,num2);
    res_wrapper=add_syscall(num1,num2);
    printf("The result of Adding %d with %d\n",num1,num2);
    printf("Result of Direct call is %d\n",res_direct);
    printf("Result of Wrapper call is %d\n",res_wrapper);

    return 0;
}
```

## 1.2: /home/user/addWrapper.h:



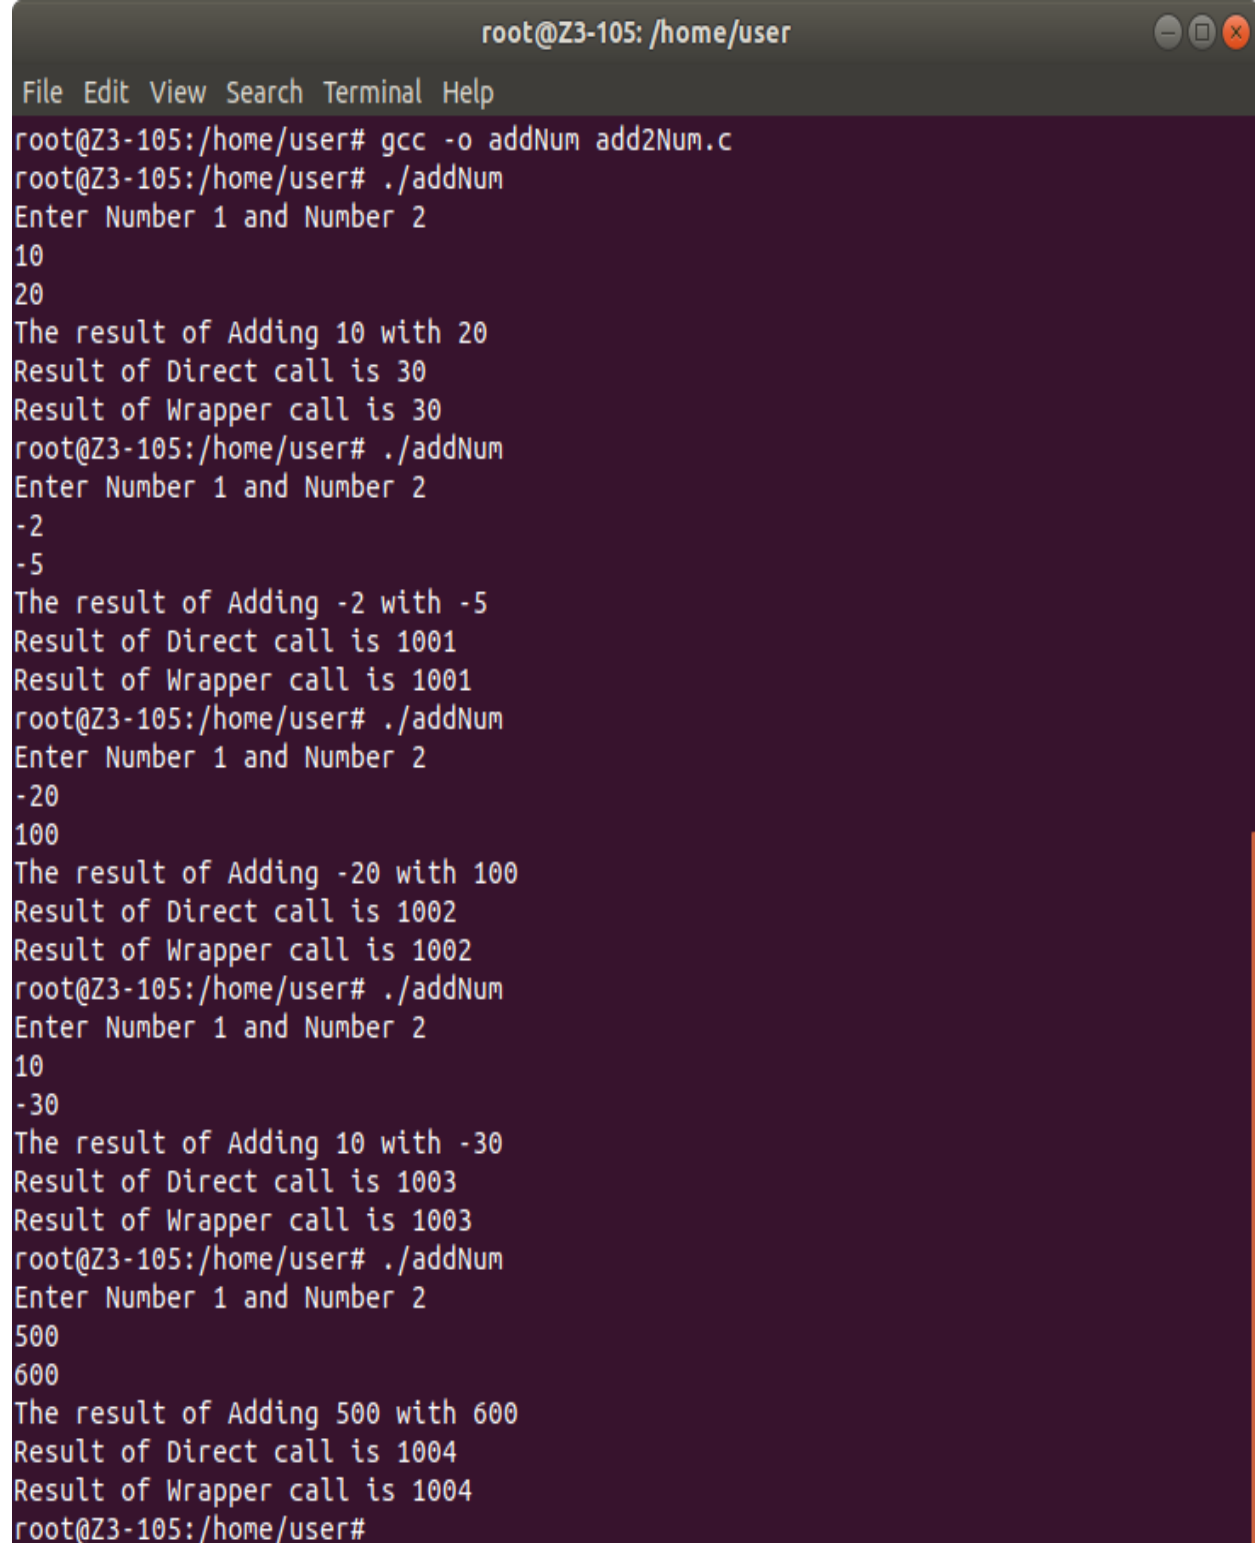
```
#ifndef __ADD_WRAPPER_H_
#define __ADD_WRAPPER_H_

#define add_syscall(num1,num2) (syscall(434,num1,num2))

#endif
```



### 1.3: Compiling and Executing the User program:



```
root@Z3-105: /home/user
File Edit View Search Terminal Help
root@Z3-105:/home/user# gcc -o addNum add2Num.c
root@Z3-105:/home/user# ./addNum
Enter Number 1 and Number 2
10
20
The result of Adding 10 with 20
Result of Direct call is 30
Result of Wrapper call is 30
root@Z3-105:/home/user# ./addNum
Enter Number 1 and Number 2
-2
-5
The result of Adding -2 with -5
Result of Direct call is 1001
Result of Wrapper call is 1001
root@Z3-105:/home/user# ./addNum
Enter Number 1 and Number 2
-20
100
The result of Adding -20 with 100
Result of Direct call is 1002
Result of Wrapper call is 1002
root@Z3-105:/home/user# ./addNum
Enter Number 1 and Number 2
10
-30
The result of Adding 10 with -30
Result of Direct call is 1003
Result of Wrapper call is 1003
root@Z3-105:/home/user# ./addNum
Enter Number 1 and Number 2
500
600
The result of Adding 500 with 600
Result of Direct call is 1004
Result of Wrapper call is 1004
root@Z3-105:/home/user#
```